

lab-06-custom-agents-intro

- [Lab 06: Introduction to Custom Copilot Agents](#)
 - [Objectives](#)
 - [Prerequisites](#)
 - [Background](#)
 - [What Are Custom Copilot Agents?](#)
 - [Mental Model: The Specialist Analogy](#)
 - [How Agents Differ From...](#)
 - [When to Use Custom Agents](#)
 - [Exercise 1: Using the Architecture Reviewer Agent \(15 minutes\)](#)
 - [Setup](#)
 - [Scenario](#)
 - [Instructions](#)
 - [Expected Outcome](#)
 - [Reflection Questions](#)
 - [Exercise 2: Compare Agent vs. Standard Chat \(10 minutes\)](#)
 - [Instructions](#)
 - [Comparison Table](#)
 - [Exercise 3: Exploring Other Agents \(5 minutes\)](#)
 - [1. Backlog Generator](#)
 - [2. Test Strategist](#)
 - [Questions](#)
 - [Key Insights](#)
 - [If Copilot Instructions Are Guardrails...](#)
 - [Agents Are Products, Not Prompts](#)
 - [Key Takeaways](#)
 - [Common Questions](#)
 - [Next Steps](#)
 - [Additional Resources](#)

Lab 06: Introduction to Custom Copilot Agents

Module: 2

Duration: 30 minutes

Part: Advanced GitHub Copilot (Part 2)

Objectives

By the end of this lab, you will:

- Understand what Custom Copilot Agents are and how they work
- Differentiate between agents, prompts, and Copilot Instructions
- Use a pre-built custom agent in Agent Mode
- Recognize when custom agents improve consistency over ad-hoc prompting

Prerequisites

- Completion of [Lab 05: Interaction Models](#)
- VS Code with GitHub Copilot extension
- Access to the TaskManager workshop repository

Background

What Are Custom Copilot Agents?

Custom Copilot Agents are **specialized AI assistants** that provide consistent, role-based guidance for specific workflows. Think of them as expert consultants you can invoke when needed.

Key Characteristics:

- **Named entities** - Selectable from the agent dropdown
- **Role-based personas** - Architecture reviewer, test strategist, backlog generator, etc.
- **Defined scope** - Clear responsibilities and constraints
- **Structured outputs** - Consistent format for results
- **Team-aligned** - Encode team practices and standards

Mental Model: The Specialist Analogy

Standard Copilot Chat = General AI Assistant
Custom Agent = Domain Expert Consultant

You wouldn't ask a general assistant to:

- Review architecture (you'd ask an architect)
- Plan testing strategy (you'd ask a QA specialist)
- Generate backlog items (you'd ask a product analyst)

Custom agents ARE those specialists.

How Agents Differ From...

Feature	Copilot Instructions	Ad-hoc Prompts	Custom Agents
Scope	Always active	One-off	Invoked on demand
Purpose	Global guardrails	Specific task	Repeatable workflow
Reusability	Implicit	Manual copy/paste	Built-in
Consistency	Background rules	Variable	Structured
Best for	Coding standards	Exploration	Workflow automation

When to Use Custom Agents

Use agents when:

- You have **repeated workflows** (reviews, planning, analysis)
- You need **consistent outputs** across team members
- You want to **encode expert knowledge** in a reusable form
- You're performing **validation or review tasks**

Don't use agents when:

- A simple prompt suffices
 - You're exploring or learning
 - The task is one-off or unique
-

Exercise 1: Using the Architecture Reviewer Agent (15 minutes)

Setup

The repository includes a pre-built **Architecture Reviewer** agent at `.github/agents/architecture-reviewer.agent.md`.

Scenario

You suspect there might be architectural issues in the Task entity or the repository implementation. You want an expert review.

Instructions

1. Locate the agent:

- Navigate to `.github/agents/architecture-reviewer.agent.md`
- Read the agent's responsibilities and constraints
- Note the structured output format

2. Open Copilot Chat and switch to Agent Mode

3. Select the Architecture Reviewer agent from the dropdown

4. Use this prompt:

Review the Task entity and TaskRepository for architectural compliance with Clean Architecture and DDD patterns.

5. Observe the agent's behavior:

- Does it follow its defined structure?
- Does it reference the layers (Domain, Application, Infrastructure, Api)?
- Does it provide the expected sections (Strengths, Concerns, Violations, Recommendations)?

6. Compare to standard Copilot Chat:

- Try the same prompt in regular Chat (without the agent)
- Note the differences in depth, structure, and consistency

Expected Outcome

The Architecture Reviewer agent should:

- Analyze the code through a Clean Architecture lens
- Identify specific boundary violations (if any)
- Provide structured findings (✅ Strengths, ⚠️ Concerns, ❌ Violations)
- Give actionable recommendations
- Reference project conventions (ADRs, DDD patterns)

Reflection Questions

1. How did the agent's response differ from standard Chat?
 2. Did the agent follow its defined output format?
 3. Would this consistency be valuable for team code reviews?
 4. What happens if you ask the agent something outside its scope?
-

Exercise 2: Compare Agent vs. Standard Chat (10 minutes)

Instructions

Perform the same task twice:

Round 1: Standard Copilot Chat

1. Open Copilot Chat (no agent selected)
2. Prompt: Review the Task domain model for DDD compliance
3. Record the response structure and depth

Round 2: Architecture Reviewer Agent

1. Switch to Agent Mode
2. Select **Architecture Reviewer** from dropdown
3. Same prompt: Review the Task domain model for DDD compliance
4. Record the response structure and depth

Comparison Table

Fill in based on your observations:

Aspect	Standard Chat	Architecture Reviewer Agent
Response Format	[Your observation]	[Your observation]
Depth of Analysis	[Your observation]	[Your observation]
Consistency	[Your observation]	[Your observation]
Actionability	[Your observation]	[Your observation]
Repeatability	[Your observation]	[Your observation]

Exercise 3: Exploring Other Agents (5 minutes)

The repository includes three custom agents. Briefly explore each:

1. Backlog Generator

- **Location:** .github/agents/backlog-generator.agent.md
- **Try:** "Generate user stories for adding task comments feature"
- **Observe:** Structured user stories with acceptance criteria

2. Test Strategist

- **Location:** .github/agents/test-strategist.agent.md
- **Try:** "Propose test scenarios for Task creation"
- **Observe:** Categorized test scenarios (unit, integration, edge cases)

Questions

- Which agent would you use most frequently in your work?
 - Can you think of other agents your team might need?
-

Key Insights

If Copilot Instructions Are Guardrails...

Copilot Instructions = Background rules always enforced
(e.g., "Use Clean Architecture, write tests first, follow DDD")

Custom Agents = Specialists you consult on demand
(e.g., "Architecture Reviewer, analyze this design")

Agents Are Products, Not Prompts

- Agents should be **versioned and reviewed** (like code)
 - Agents **encode team knowledge** and standards
 - Agents provide **repeatable, consistent outcomes**
 - Agents improve **onboarding** (new team members use the same expert guidance)
-

Key Takeaways

- ✓ Custom agents provide role-based expertise on demand
 - ✓ Agents ensure consistency across team members
 - ✓ Agents are reusable - define once, use repeatedly
 - ✓ Agents complement Instructions - not a replacement
 - ⚠ Agents require maintenance - treat them as team assets
-

Common Questions

Q: Can I use multiple agents in one session?

A: Yes! Switch agents as needed for different workflow steps.

Q: Do agents replace Copilot Instructions?

A: No. Instructions are always-on guardrails; agents are on-demand specialists.

Q: Can I create my own agent?

A: Absolutely! That's covered in [Lab 08: Agent Design](#) and [Lab 09: Build Your Own](#).

Q: What if an agent gives incorrect advice?

A: Agents are assistants, not authorities. You're accountable for the final decision. Iterate on agent instructions to improve accuracy.

Next Steps

In [Lab 07: Workflow Agents in Action](#), you'll apply these agents to real development workflows and compare their outputs to ad-hoc prompting.

Additional Resources

- [Custom Agent Catalog](#)
- [Agent vs Instructions vs Prompts Diagram](#)
- [GitHub Documentation: Custom Agents](#)